

Lecture 5: GPU programming

Valentin Churavy (@vchuravy)



Levels of Parallelism in Hardware

1. Instruction-Level Parallelism
2. Data-Level Parallelism
 - a. SIMD/Vector
 - b. GPUs
3. Thread-Level Parallelism



Instruction-Level Parallelism

Assembly-esque Julia syntax

```
function f(A, x)
    i = length(A)

    @label Loop
    a = A[i]      # Load
    c = a + x     # Add
    A[i] = c      # Store
    i = i - 1     # Decrement
    i > 0 && @goto Loop

    return A
end
```

RISC-V Assembly

```
...
Loop: fld      f0, 0(x1)
      fadd.d   f4, f0, f2
      fsd      f4, 0(x1)
      addi     x1, x1, -8
      bnez     x1, Loop
...
```

What are the data-dependencies in the loop?

Pipeline Scheduling: Latency

Naive scheduling

```
@label Loop
a = A[i]           # Cycle 1
# Stall           # Cycle 2
c = a + x          # Cycle 3
# Stall           # Cycle 4
# Stall           # Cycle 5
A[i] = c           # Cycle 6
i = i - 1          # Cycle 7
i > 0 && @goto Loop # Cycle 8
```

- Load latency: 1 cycle
- Float arithmetic latency: 2 cycle
- Integer arithmetic latency: 0 cycle

After reordering

```
@label Loop
a = A[i]           # Cycle 1
i = i - 1          # Cycle 2
c = a + x          # Cycle 3
# Stall           # Cycle 4
# Stall           # Cycle 5
A[i+1] = c         # Cycle 6
i > 0 && @goto Loop # Cycle 7
```

- How many cycles are overhead: 2
- How many cycles are stalls: 2
- How many cycles are actually work: 3

Note: `A[i+1]` is free since it can be precomputed relative to `A[i]`

Unrolling

Goal: Reduce overhead relative to work

- Replicate loop body
 - Rename variables
 - Change stride and exit condition
 - Add tail loop that deals with non-multiples of the unroll factor
-
- Do we still have stalls?: Yes
 - How many stall cycles: 12
 - How many cycles are overhead: 2
 - How many cycles are actually work: 12

Unroll factor: 4

```
@label Loop
a = A[i]
c = a + x
A[i] = c
a1 = A[i-1]
c1 = a1 + x
A[i-1] = c1
a2 = A[i-2]
c2 = a2 + x
A[i-2] = c2
a3 = A[i-3]
c3 = a3 + x
A[i-3] = c3
i = i - 4
i > 4 && @goto Loop
```

Note: `A[i-1]` is free since it can be precomputed relative to `A[i]`

Loop unrolling and re-ordering

Loop with stalls annotated

```
@label Loop
a = A[i]
# Stall
c = a + x
# Stall
# Stall
A[i] = c
a1 = A[i-1]
# Stall
c1 = a1 + x
# Stall
# Stall
A[i-1] = c1
a2 = A[i-2]
# Stall
c2 = a2 + x
# Stall
# Stall
A[i-2] = c2
a3 = A[i-3]
# Stall
c3 = a3 + x
# Stall
# Stall
A[i-3] = c3
i = i - 4
i > 4 && @goto Loop
```

Re-ordered loop

```
@label Loop
a = A[i]
a1 = A[i-1]
a2 = A[i-2]
a3 = A[i-3]
c = a + x
c1 = a1 + x
c2 = a2 + x
c3 = a3 + x
A[i] = c
A[i-1] = c1
A[i-2] = c2
A[i-3] = c3
i = i - 4
i > 4 && @goto Loop
```

- How many stalls? 0
- How many overhead cycles: 2
- How many cycles are actually work: 12

Instruction-Level Parallelism

What has your processor done for you recently?

As programmers we have the mental model that a processor takes your instructions in **linear** order

In reality they exploit instruction-level parallelism and have deep cache hierarchies

Ways to exploit instruction-level parallelism:

- Superscalar / Multiple issue
- Out-of-order execution / Speculation
- Prediction

Data-parallelism

Instead of writing serial code and expect the processor to pick up our slack, we could also write explicit parallel code, this can be especially beneficial for inherently data-parallel codes

- Explicit Vector Programming
- GPU Programming

What are data-parallel programs?

- ``map`/`broadcast``
- Matrix multiply

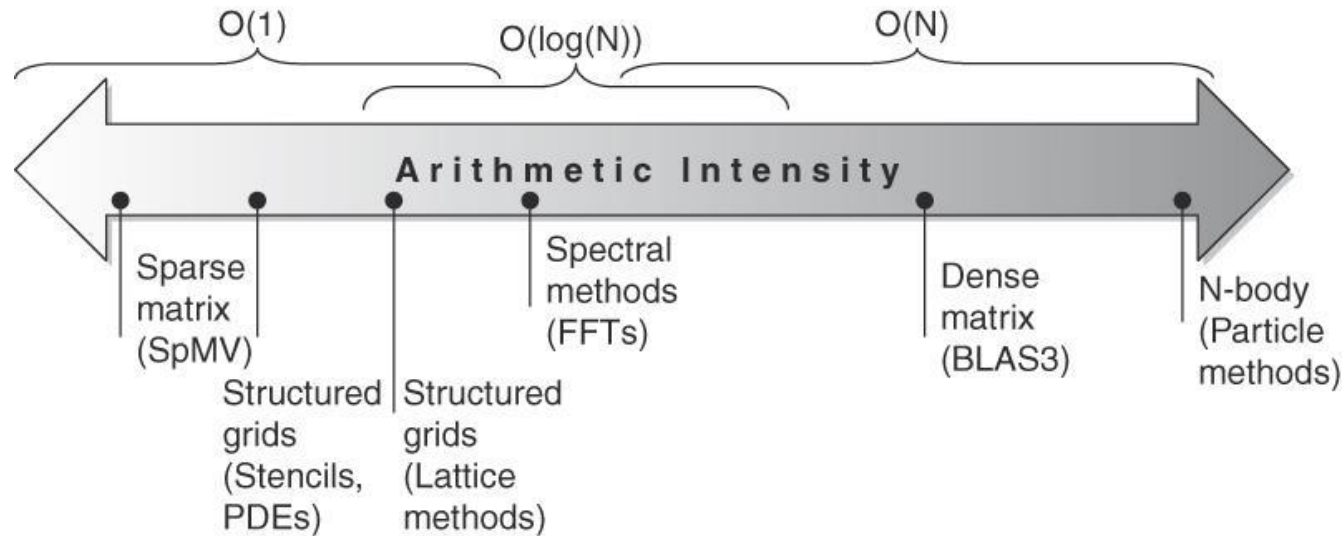


Figure 4.10 Arithmetic intensity, specified as the number of floating-point operations to run the program divided by the number of bytes accessed in main memory [Williams et al. 2009]. Some kernels have an arithmetic intensity that scales with problem size, such as dense matrix, but there are many kernels with arithmetic intensities independent of problem size.

SIMD (Explicitly vectorized)

```
using SIMD
A = rand(Float64, 64)
T = Vec{4, Float64}
x = 1.0

for i in 1:4:length(A)
    a = vload(T, A, i)
    c = a + x
    vstore(c, A, i)
end
```

- Stalls are only per instruction, and not per element
- Reduced overhead
 - 3 instructions processing 4 elements
 - 2 overhead instructions
 - Overhead is amortized across 4 elements.

- We can remove stalls similar to what we did for the serial code
 - Pipelining
 - Interleaving and Unrolling
- Latencies will be higher

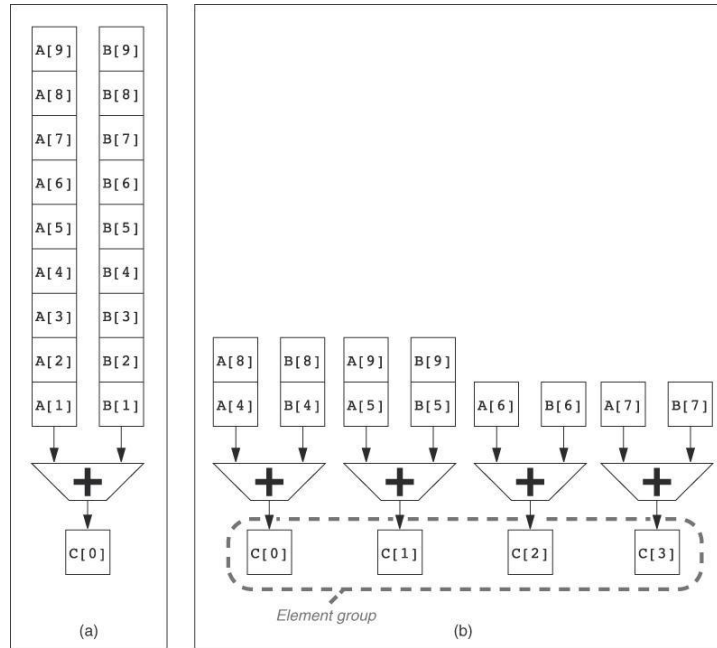


Figure 4.4 Using multiple functional units to improve the performance of a single vector add instruction, $C = A + B$.

The vector processor (a) on the left has a single add pipeline and can complete one addition per cycle. The vector processor (b) on the right has four add pipelines and can complete four additions per cycle. The elements within a single vector add instruction are interleaved across the four pipelines. The set of elements that move through the pipelines together is termed an *element group*. (Reproduced with permission from Asanovic [1998].)

Vector gather and scatter

`vload` loads a block of memory.

How do we work with sparse data and indirect indexing?

- Index vector of the same SIMD length
- Gather to load data
- Scatter to write data

```
idx = Vec((1, 3, 2, 4))  
v = vgather(A, idx)  
vscatter(v, A, idx)
```

Less efficient than vload

How do we handle code that branches

```
A = rand(Int64, 64)
for i in 1:length(A)
    a = A[i]
    if a % 2 == 0
        A[i] = -a
    end
end
```

Data or index dependent controlflow

```
using SIMD
A = rand{Int64, 64}
T = Vec{4, Int64}

for i in 1:4:length(A)
    a = vload(T, A, i)
    mask = a % 2 == 0           # calculate mask
    b = -a                     # If branch
    c = vifelse(mask, b, a)    # merge results
    vstore(c, A, i)
end
```

Hardware solution:

- Vector predication (NEC Aurora VX)
- Masked load and store (SIMD Intel CPU)

Compiler solution:

- ISPC
- OpenCL/CUDA

GPU (implicit vectorization)

```
pkg> add CUDAnative
julia> using CUDAnative
julia> function say(num)
    @cuprintf("Thread %ld says: %ld\n",
              threadIdx().x, num)
    return
end
julia> @cuda threads=4 say(42)
Thread 1 says: 42
Thread 2 says: 42
Thread 3 says: 42
Thread 4 says: 42
```

- Lane index:
 `threadIdx().x`
- Number of lanes:
 `threads=4`
- Call of a kernel:
 `@cuda`
- Native vector width:
 32

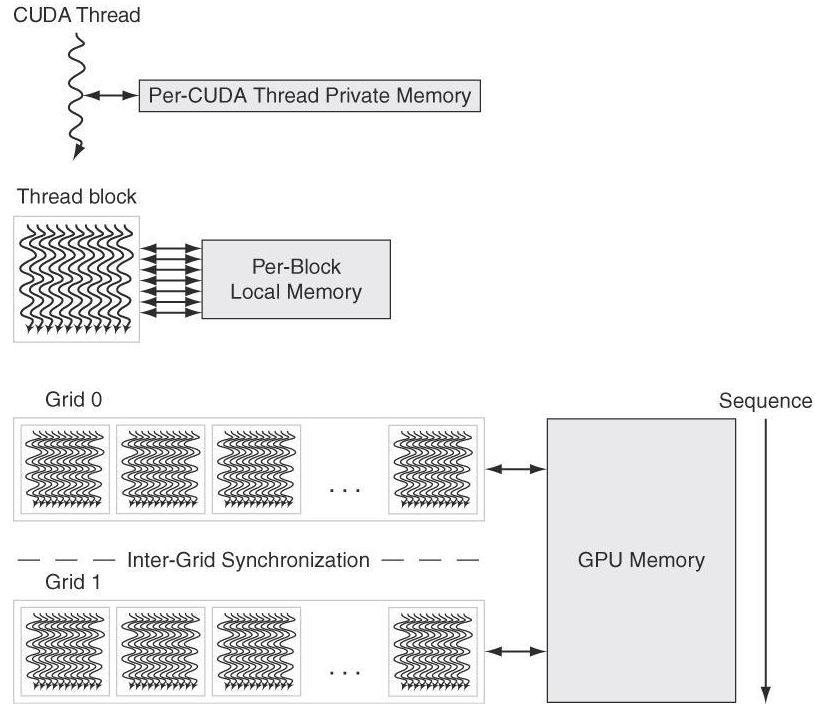


Figure 4.18 GPU Memory structures. GPU Memory is shared by all Grids (vectorized loops), Local Memory is shared by all threads of SIMD instructions within a thread block (body of a vectorized loop), and Private Memory is private to a single CUDA Thread.

Multiple layers of parallelism

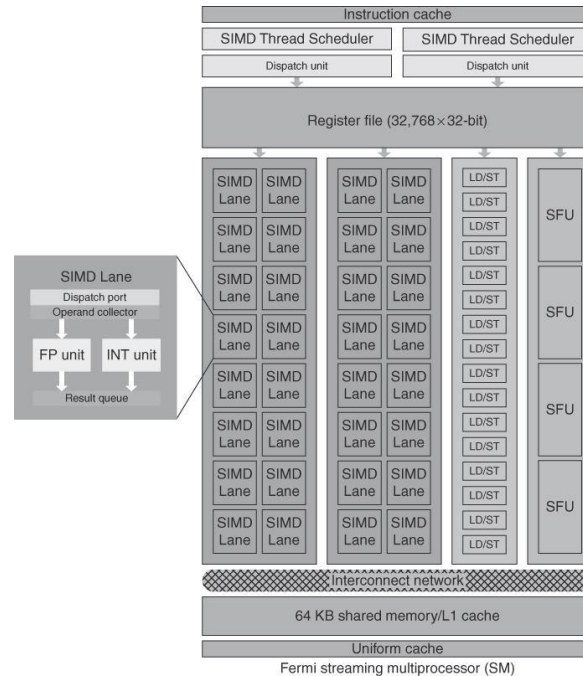


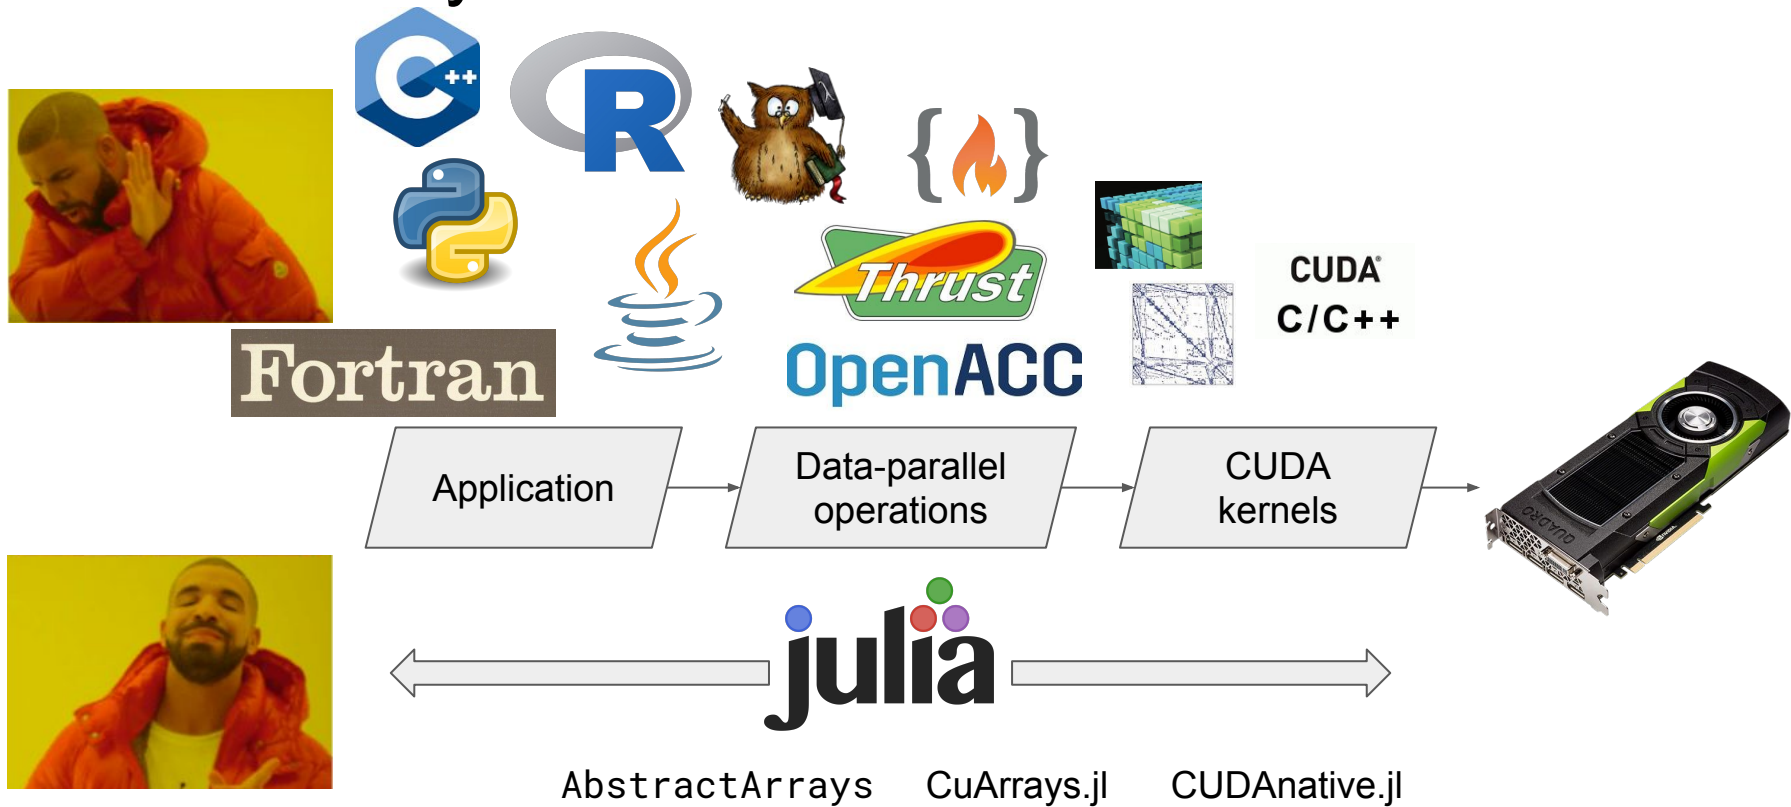
Figure 4.20 Block diagram of the multithreaded SIMD Processor of a Fermi GPU. Each SIMD Lane has a pipelined floating-point unit, a pipelined integer unit, some logic for dispatching instructions and operands to these units, and a queue for holding results. The four Special Function units (SFUs) calculate functions such as square roots, reciprocals, sines, and cosines.

GPU programming in Julia

With material from Tim Besard




How to train your GPU: 10.000 foot view



Why should you care?

1) Performance

2) Powerful abstractions



Show me what you got

```
julia> function say(f)
    i = threadIdx().x
    @cuprintf("Thread %ld says: %ld\n",
              i, f(i))
    return
end
```

```
julia> @cuda say(x->x+1)
Thread 1 says: 2
```

Show me what you got

```
julia> a = CuArray([1., 2., 3.])

julia> function apply(op, a)
    i = threadIdx().x
    a[i] = op(a[i])
    return
end

julia> @cuda threads=length(a) map(x->x^2, a)

julia> a
3-element CuArray{Float32,1}:
 1.0
 4.0
 9.0
```

```
julia> @device_code_ptx @cuda apply(x->x^2, a)
apply(.param .b8 a[16])
{
    ld.param.u64    %rd1, [a+8];
    mov.u32         %r1, %tid.x;

    // index calculation
    mul.wide.u32     %rd2, %r1, 4;
    add.s64          %rd3, %rd1, %rd2;
    cvta.to.global.u64 %rd4, %rd3;

    ld.global.f32    %f1, [%rd4];
    mul.f32          %f2, %f1, %f1;
    st.global.f32    [%rd4], %f2;

    ret;
}
```

Julia array abstractions for high-level GPU programming

551 commits

10 branches

16 releases

23 contributors

View license

No GPU programming experience

Data-parallel programming model



Not just another array library

```
julia> a = CuArray([1,2,3])  
3-element CuArray{Int64,1}:
```

```
1  
2  
3
```

dot syntax



```
julia> function apply(op, a)  
    i = threadIdx().x  
    a[i] = op(a[i])  
end  
julia> @cuda threads=length(a) apply(op, a)  
  
julia> map(op, a)  
  
julia> reduce(binop, a)  
6  
  
julia> broadcast(+, [1], [2 2], [3 3; 3 3])  
2×2 CuArray{Int64,2}:  
6 6  
6 6  
  
julia> [1] .+ [2 2] .+ [3 3; 3 3]  
2×2 CuArray{Int64,2}:  
6 6  
6 6
```


Not just another array library

```
julia> a = CuArray([1f0, 2f0, 3f0])
```

```
3-element CuArray{Float32,1}:
```

```
1.0
```

```
2.0
```

```
3.0
```

```
julia> f(x) = 3x^2 + 5x + 2
```


```
julia> a .= f.(2 .* a.^2 .+ 6 .* a.^3 .- sqrt.(a))
```

```
3-element CuArray{Float32,1}:
```

```
184.0
```

```
9213.753
```

```
96231.72
```



Single kernel!

- Fully specialized
- Highly optimized
- Great performance



Vendor libraries

```
julia> a = CuArray{Float32}(undef, (2,2));
```

CURAND

```
julia> rand!(a)
2×2 CuArray{Float32,2}:
0.73055  0.843176
0.939997 0.61159
```

CUBLAS

```
julia> a * a
2×2 CuArray{Float32,2}:
1.32629  1.13166
1.26161  1.16663
```

CUSOLVER

```
julia> LinearAlgebra.qr!(a)
CuQR{Float32,CuArray{Float32,2}}
with factors Q and R:
Float32[-0.613648 -0.78958; -0.78958 0.613648]
Float32[-1.1905 -1.00031; 0.0 -0.290454]
```

CUFFT

```
julia> CUFFT.plan_fft(a) * a
2-element CuArray{Complex{Float32},1}:
-1.99196+0.0im  0.589576+0.0im
-2.38968+0.0im -0.969958+0.0im
```

CUDNN

```
julia> softmax(real(ans))
2×2 CuArray{Float32,2}:
0.15712  0.32963
0.84288  0.67037
```

CUSPARSE

```
julia> sparse(a)
2×2 CuSparseMatrixCSR{Float32,Int32}
with 4 stored entries:
 [1, 1] = -1.1905
 [2, 1] = 0.489313
 [1, 2] = -1.00031
 [2, 2] = -0.290454
```

Effective GPU Programming

How do you *actually* use this stuff?

Types and gradients, including Forward.gradient

■ <https://discourse.julialang.org/t/types-and-gradients-including-forward-gradient/946>

```
using LinearAlgebra
```


```
loss(w,b,x,y) = sum(abs2, y - (w*x .+ b)) / size(y,2)
loss∇w(w, b, x, y) = ...
lossdb(w, b, x, y) = ...
```

```
function train(w, b, x, y ; lr=.1)
    w -= lmul!(lr, loss∇w(w, b, x, y))
    b -= lr * lossdb(w, b, x, y)
    return w, b
end
```

```
n = 100
p = 10
x = randn(n,p)'
y = sum(x[1:5,:]; dims=1) .+ randn(n)'*0.1
w = 0.0001*randn(1,p)
b = 0.0
```

```
for i=1:50
    w, b = train(w, b, x, y)
end
```

```
x = CuArray(x)
y = CuArray(y)
w = CuArray(w)
```



Fin.

cuArrays vs CUDANative

■ <https://discourse.julialang.org/t/cuarrays-vs-cudanative/17504>

```
function diff_y(a, b)
    s = size(a)
    for j = 1:s[2]
        for i = 1:s[1]
            @inbounds a[i,j] = b[i,j+1] - b[i,j]
        end
    end
end

N = 64
nx = N^2
ny = N
a = ones(Float32, nx, ny-1)
b = ones(Float32, nx, ny)

julia> using BenchmarkTools
julia> @btime diff_y($a,$b);
39.599 μs (0 allocations: 0 bytes)

julia> @btime diff_y($(CuArray(a)),$(CuArray(b)));
4.499 s (3354624 allocations: 165.38 MiB)
```

Performance killers

1. Scalar iteration is sloooooow

```
function diff_y(a, b)
    s = size(a)
    for j = 1:s[2]
        for i = 1:s[1]
            @inbounds a[i,j] = b[i,j+1] - b[i,j]
        end
    end
end
```

```
julia> CuArrays.allowscalar(false)

julia> diff_y(CuArray(a), CuArray(b))
ERROR: scalar getindex is disallowed
Stacktrace:
...
[5] getindex at ./abstractarray.jl
[6] diff_y(::CuArray, ::CuArray)
    at ./REPL[109]:5
...
```

```
function diff_y(a, b)
    s = size(a)
    for j = 1:s[2]
        @inbounds a[:,j] .= b[:,j+1] - b[:,j]
    end
end

julia> @btime diff_y($(CuArray(a)), $(CuArray(b)));
2.503 ms (16884 allocations: 661.50 KiB)
```

Performance killers

2. Avoid multiple kernels

```
function diff_y(a, b)
```

```
    a .= @views b[:, 2:end] .- b[:, 1:end-1]
```

```
end
```

```
julia> @btime diff_y($(CuArray(a)),$(CuArray(b))); 39.057 μs (40 allocations: 2.08 KiB)
```

```
julia> @btime diff_y($a,$b);  
39.599 μs (0 allocations: 0 bytes)
```

```
function diff_y(a, b)
```

```
    s = size(a)
```

```
    for j = 1:s[2]
```

```
        @inbounds a[:,j] .= b[:,j+1] - b[:,j]
```

```
    end
```

```
end
```

```
julia> @btime diff_y($(CuArray(a)),$(CuArray(b))); 2.503 ms (16884 allocations: 661.50 KiB)
```

Performance killers

3. Bump the problem size

```
julia> N = 256
```

```
julia> @btime diff_y($(CuArray(a)),$(CuArray(b)));  
1.494 ms (40 allocations: 2.08 KiB)
```

```
julia> @btime diff_y($a,$b);  
11.719 ms (2 allocations: 128 bytes)
```

4. Keep data on the GPU

```
julia> @btime diff_y(CuArray($a),CuArray($b));  
72.050 ms (93 allocations: 255.50 MiB)
```


Strengths

1. Single, productive programming language
2. Platform-independent, generic code
3. High-level, zero-cost abstractions
4. Great performance potential
5. **Composability**
6. **Optimizability**

Composability

Separation of concerns

```
julia> map(x->x^2, CuArray([1 2 3]))
```

- *what* is computed
- *where* does it happen
- *how* is it computed

CUDAnative.jl	2383 LOC
GPUArrays.jl	1468 LOC
CuArrays.jl	859 LOC (without libraries)

Composability: reuse of libraries

```
loss(w,b,x,y) = sum(abs2, y - (w*x .+ b)) / size(y,2)
```

```
julia> loss(w,b,x,y)
4.222961132618434
```

```
julia> loss∇w(w, b, x, y)
1×10 CuArray{Float64,2}:
-1.365  -1.961  -1.14  -2.023  -1.981  -0.2993  -0.2667  -0.07669  -1.038  -0.1823
```

```
using ForwardDiff
```

```
loss∇w(w, b, x, y) = ForwardDiff.gradient(w -> loss(w, b, x, y), w)
```

```
julia> @which mul!(w, w, x)
mul!(...) in CuArrays.CUBLAS at src/blas/highlevel.jl
```

```
julia> @which mul!(w, w, ForwardDiff.Dual.(x))
mul!(...) in CuArrays at src/generic_matmul.jl
```

Optimizability: it's



the way down

1. Rewrite using array abstractions
using `CuArrays` + generic code
2. Avoid GPU antipatterns
3. Specialize with broadcast expressions
4. Specialize with GPU kernels

Use the Tools

Tools

1. Reflection and introspection

```
julia> using CUDAnative
```

```
julia> @device_code_llvm curand(2) .+ 2
```

```
define void @ptxcall_anonymous(...) {  
    ...  
}
```

```
@device_code_{lowered,typed,warntype,llvm,ptx,sass}
```

```
julia> ENV["JULIA_DEBUG"] = "CUDAnative"
```

```
julia> curand(2) .+ 2;
```

```
└─ Debug: Compiled getField(GPUArrays, ...>() to PTX 3.5.0 for SM 3.5.0 using 8 registers.  
└─ Memory usage: 0 bytes local, 0 bytes shared, 0 bytes constant  
└─ @ CUDAnative ~/Julia/CUDAnative/src/execution.jl
```

Tools

2. Performance measurements

```
julia> const x = CuArray{Float32}(undef, 1024)

julia> using BenchmarkTools
julia> @benchmark CuArrays.@sync(identity.($x))
BenchmarkTools.Trial:
  memory estimate:  1.34 KiB
  allocs estimate:  33
  -----
  minimum time:     13.824 μs (0.00% GC)
  median time:      16.361 μs (0.00% GC)
  mean time:        16.489 μs (0.00% GC)
  maximum time:     401.689 μs (0.00% GC)
  -----
  samples:          10000
  evals/sample:     1
```

Tools

2. Performance measurements

```
julia> const x = CuArray{Float32}(undef, 1024)

julia> using BenchmarkTools
julia> @benchmark CuArrays.@sync(identity.($x))
BenchmarkTools.Trial:
  minimum time:      13.824 μs (0.00% GC)
  maximum time:      401.689 μs (0.00% GC)

julia> CuArrays.@time CuArrays.@sync identity.(x);
0.000378 seconds (57 CPU allocations: 1.938 KiB)
(1 GPU allocation: 4.000 KiB)
```

Tools

2. Performance measurements

```
julia> const x = CuArray{Float32}(undef, 1024)
```

```
julia> using BenchmarkTools
```

```
julia> @benchmark CuArrays.@sync(identity.($x))
```

```
BenchmarkTools.Trial:
```

```
  minimum time:      13.824 μs (0.00% GC)
```

```
  maximum time:      401.689 μs (0.00% GC)
```

```
julia> CuArrays.@time CuArrays.@sync identity.(x);  
0.000378 seconds (57 CPU allocations: 1.938 KiB)  
          (1 GPU allocation: 4.000 KiB)
```

```
julia> using CUDAdrv
```

```
julia> CUDAdrv.@elapsed identity.(x)
```

```
5.888f-6
```

**Accurate measurements of
possible short-running code**

Memory allocation behavior

Application performance metrics

Tools

3. Profiling

```
$ nvprof --profile-from-start off julia
```

```
julia> const x = CuArray{Float32}(undef, 1024)
julia> identity.(x)
```

```
julia> CUDAdrv.@profile begin
        identity.(x)
    end
```

```
julia> exit()
```

```
==22272== Profiling result:
```

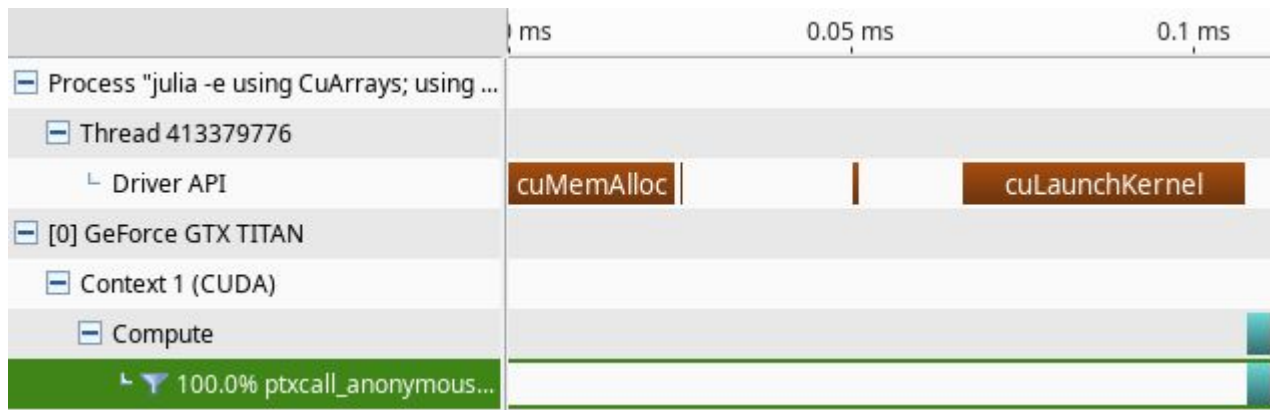
	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	100.00%	3.5520us	1	3.5520us	3.5520us	3.5520us	3.5520us	ptxcall_anonymous
API calls:	61.70%	39.212us	1	39.212us	39.212us	39.212us	39.212us	cuLaunchKernel
	37.36%	23.745us	1	23.745us	23.745us	23.745us	23.745us	cuMemAlloc
	0.93%	592ns	2	296ns	222ns	370ns		cuCtxGetCurrent

Tools

3. Profiling

```
$ nvvp julia
```

```
julia> identity.(CuArray{Float32}(undef, 1024))
```

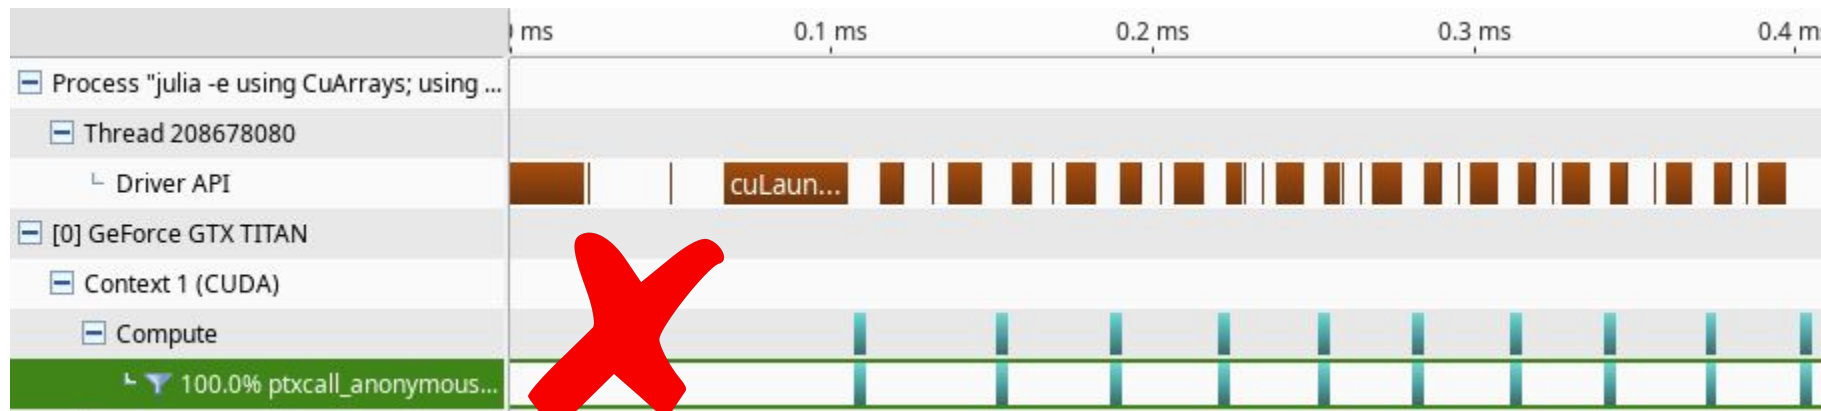


Tools

3. Profiling

```
$ nvvp julia
```

```
julia> identity.(CuArray{Float32}(undef, 1024))
```




Tools

3. Profiling

```
$ nvvp julia
```

```
julia> sin.(CuArray{Float32}(undef, 1024, 1024))
```



Effectively using GPUs with Julia

Tim Besard (@maleadt)

<http://julialang.slack.com/>
<https://discourse.julialang.org/c/domain/gpu>

